

GetResilience v1.3

System Configuration Document

Onno Hansen-Staszyński, 2026

1. Overview

This document describes the system configuration for **GetResilience**, a citizen-led platform for establishing grassroots resilience councils. GetResilience enables communities to detect, report, and respond to information manipulation through a structured, legally compliant process.

The platform is built around the principle that citizens can play an active role in defending the information environment, while operating within the boundaries of the Digital Services Act (DSA) and the General Data Protection Regulation (GDPR).

The configuration described here is version 1.3 (Hardened), which introduces stricter validation, role-based access, and enhanced audit capabilities.

2. System Identity

Each deployment of GetResilience must identify itself with the following information:

Field	Description	Requirements
Name	The name of the system deployment	Minimum 3 characters
Version	Semantic version number (e.g. 1.3 or 1.4)	Must follow major.minor or major.minor.patch format
Legal name	The legal name of the operating entity	Minimum 3 characters
Jurisdiction	Country or legal jurisdiction	Minimum 2 characters (e.g. country code)
Roles	The roles the entity fulfils	One or more of: Data Controller, Hosting Provider

3. Dual-Track Architecture

GetResilience operates on a **dual-track model**, separating legally actionable content from broader pattern analysis. This ensures that each type of signal is handled with the appropriate legal basis and level of attribution.

3.1 Track A — Illegal Content Notices

Track A handles content that is potentially **illegal under applicable law**. Reports in this track result in formal legal notices that are attributable to the reporting council.

Aspect	Detail
Scope	Illegal content only
Legal basis (GDPR Art. 6)	Legitimate interest or legal obligation
Legal basis (GDPR Art. 9)	Legal claims (required for this track)
Output type	Legal notice
Attributable	Yes — the output is attributed to the council

3.2 Track B — Pattern Analysis & Intelligence

Track B covers content and patterns that may be **harmful but not illegal**, including foreign information manipulation and interference (FIMI) signals and cross-platform patterns. The output is aggregated intelligence that is **not attributable** to any individual reporter.

Aspect	Detail
Scope	Harmful but legal content, FIMI signals, cross-platform patterns
Legal basis (GDPR Art. 6)	Legitimate interest or legal obligation
Legal basis (GDPR Art. 9)	Substantial public interest (required for this track)
Output type	Aggregated intelligence
Attributable	No — the output cannot be traced to individual reporters

4. Actor Model & Permissions

The system uses a **layered actor model** that defines who can do what within the platform. Each layer has specific roles, permissions, and requirements.

Layer	Actor	Permissions / Requirements
L0 — Public	Anonymous user	Can submit signals (reports). No account or training required.
L2 — Council Member	Resilience Council member	Must hold an EMOD certification. Can participate in verification and review.
L3 — Consensus	Council (collective)	Decisions require a 3-of-n consensus — at least 3 members must agree.

Note: Layer 1 (L1) is reserved for future use and is not defined in this version of the configuration.

5. Signal Intake

When a user submits a signal (report) to the system, the following information is collected:

- **URL** — the web address of the content being reported (must be a valid URI)
- **Platform** — the platform where the content was found (e.g. Facebook, X, Telegram)
- **Classification** — optional further categorisation of the content

All incoming signals are classified as **non-attributable** by default. This means the identity of the person who submitted the signal is not linked to the report, protecting the privacy of reporters.

6. Workflows

6.1 Track A Workflow

Reports in Track A follow a structured six-step process:

- **Step 1: Signal submission** — A user submits a report about potentially illegal content.
- **Step 2: Verification** — The report is verified for completeness and relevance.
- **Step 3: Draft notice** — A formal legal notice is drafted based on the verified report.
- **Step 4: Review** — Council members review the draft notice.
- **Step 5: Consensus** — At least 3 council members must agree on the notice (3-of-n rule).
- **Step 6: Submission** — The finalised notice is submitted to the relevant platform or authority.

6.2 Track B Workflow

Track B uses a five-step analytical pipeline:

- **Step 1: Collection** — Signals and data points are gathered from multiple sources.
- **Step 2: Clustering** — Related signals are grouped together to identify patterns.
- **Step 3: Pattern detection** — Algorithms and analysts identify coordinated or recurring patterns.
- **Step 4: Aggregation** — Findings are aggregated into anonymised intelligence reports.
- **Step 5: Output** — The aggregated intelligence is published or shared with stakeholders.

6.3 Article 18 Notifications

Under Article 18 of the DSA, certain notifications must be sent to affected parties. The system supports configurable notification templates for this purpose.

7. Legal Compliance

7.1 Digital Services Act (DSA)

The system is designed to operate within the **safe harbour** provisions of the DSA. This means:

- **No active role** — the platform does not actively seek out illegal content; it processes reports submitted by users.
- **Knowledge condition** — the platform only acts on content after it has been validated through the council process (post-validation only).
- **Expeditious action** — once content is validated as illegal, the platform acts promptly to issue notices.

7.2 General Data Protection Regulation (GDPR)

The system operates under GDPR with the following configuration:

- **Data controller** — the operating entity (identified in Section 2) acts as the data controller.
- **Legal basis** — processing is based on legitimate interest or legal obligation (Art. 6), with additional grounds under Art. 9 where special categories of data are involved.

8. Security

The system includes the following security measures:

Measure	Description
Rate limiting	Enabled to prevent abuse and ensure fair use of the platform.
CAPTCHA	Enabled to prevent automated or bot-driven submissions.
Data classification	All data is classified into one of four levels: Public, Internal, Restricted, or Sensitive. Access controls are applied accordingly.

9. Audit & Logging

The system maintains a comprehensive audit log to ensure transparency and accountability. Audit logging is **enabled by default** and cannot be disabled in production environments.

Each audit log entry contains:

- **Event ID** — a unique identifier for the event (UUID format)
- **Timestamp** — the date and time of the event (ISO 8601 format)
- **Actor** — who performed the action
- **Action** — what was done
- **Trace ID** — optional, for linking related events across the system

10. Data Classification Levels

Level	Description
Public	Information that can be freely shared. No access restrictions.
Internal	Information intended for internal use within the resilience council.
Restricted	Sensitive information with limited access. Requires authorisation.
Sensitive	Highly sensitive data (e.g. personal data of reporters). Strictest access controls apply.

Appendix: Version History

Version	Date	Changes
1.0	2025	Initial system configuration
1.3	2026	Hardened configuration: stricter validation, dual-track architecture, enhanced audit logging, GDPR/DSA compliance framework